FALL 2024 COS597R:

DEEP DIVE INTO LARGE LANGUAGE MODELS

Danqi Chen, Sanjeev Arora

PLI PRINCETON UNIVERSITY

Lecture 9: Alignment —What, Why, How

https://princeton-cos597r.github.io/

Preference optimization (some comments)

Rewards, Preferences, Chess, etc.

$$\Pr(i > j) = \frac{e^{\beta_i}}{e^{\beta_i} + e^{\beta_j}}$$

Bradley-Terry Model of preferences

β_i = ``quality" of i

Given a set of observed preferences, can fit β_i's

What is max-likelihood β_i's given observed outcomes?

ELO ratings (chess): Given win-loss history over time, can estimate scalar rating (β_i's) for all players ("ELO Rating" = $400\beta_i$)

Rank* (UB)	Model	Arena Score
1	o1-preview	1339
1	ChatGPT-4o-latest (2024-09-03)	1337
3	o1-mini	1314
4	Gemini-1.5-Pro-Exp-0827	1299
4	Grok-2-08-13	1293
6	GPT-4o-2024-05-13	1285
7	GPT-4o-mini-2024-07-18	1272
7	Claude 3.5 Sonnet	1269

Meaning of Learning Objectives

P: teacher $\quad Q$: learner

$$KL(P||Q) = E_{y\sim P}[\log \frac{P(y)}{Q(y)}] \quad \text{vs} \quad KL(Q||P) = E_{y\sim Q}[\log \frac{Q(y)}{P(y)}]$$

"Forward KL" "Reverse KL"

Discuss:

1. What do these objectives mean, and what training scenarios do they correspond to?

2. If teacher gives low/high probability to some y's, how does this shape Q ?

3. If student gives probability almost 0 to some y's how does this shape Q

(Note: In alignment we want student to give zero (very low) probability to some y's

Two behaviors

$$KL(P||Q) = E_{y\sim P}[\log \frac{P(y)}{Q(y)}]$$

$$KL(Q||P) = E_{y\sim Q}[\log \frac{Q(y)}{P(y)}]$$

Mode-covering

Q gives high-ish probability to y's where $P(y)$ is high; free to do anything for y's where $P(y)$ is low

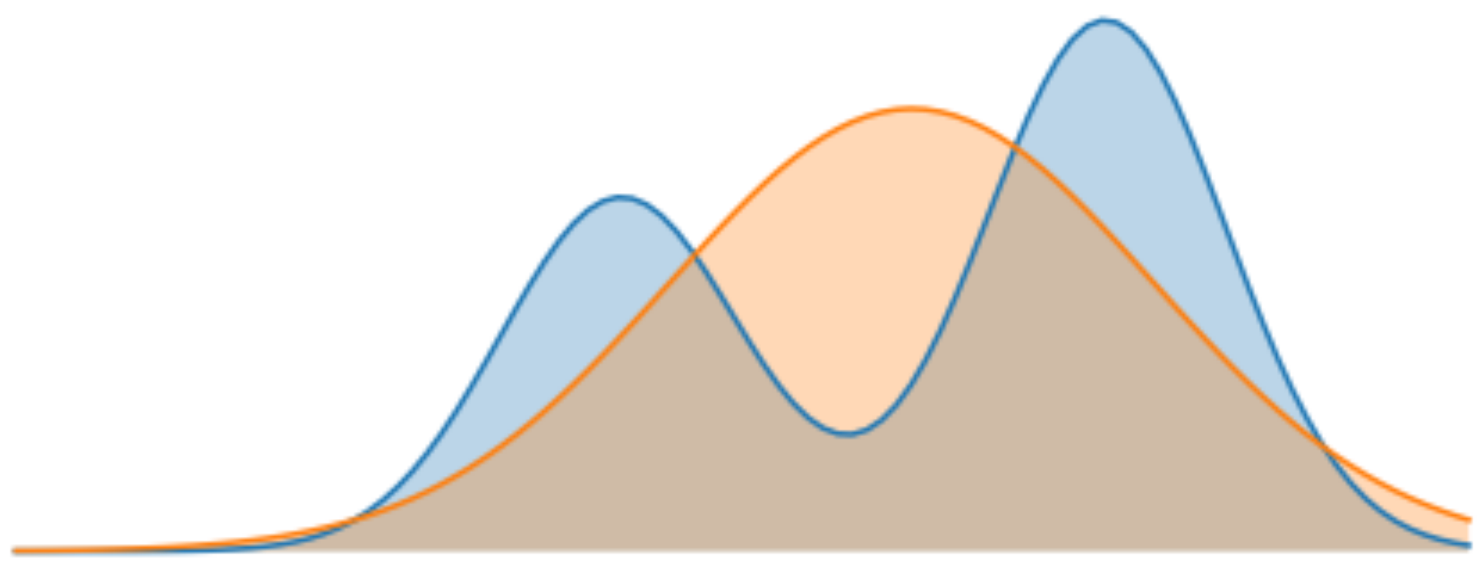

Mode-seeking

Q gives high-ish probability only to y's where $P(y)$ is high.

Give low probability to y where $P(y)$ is low

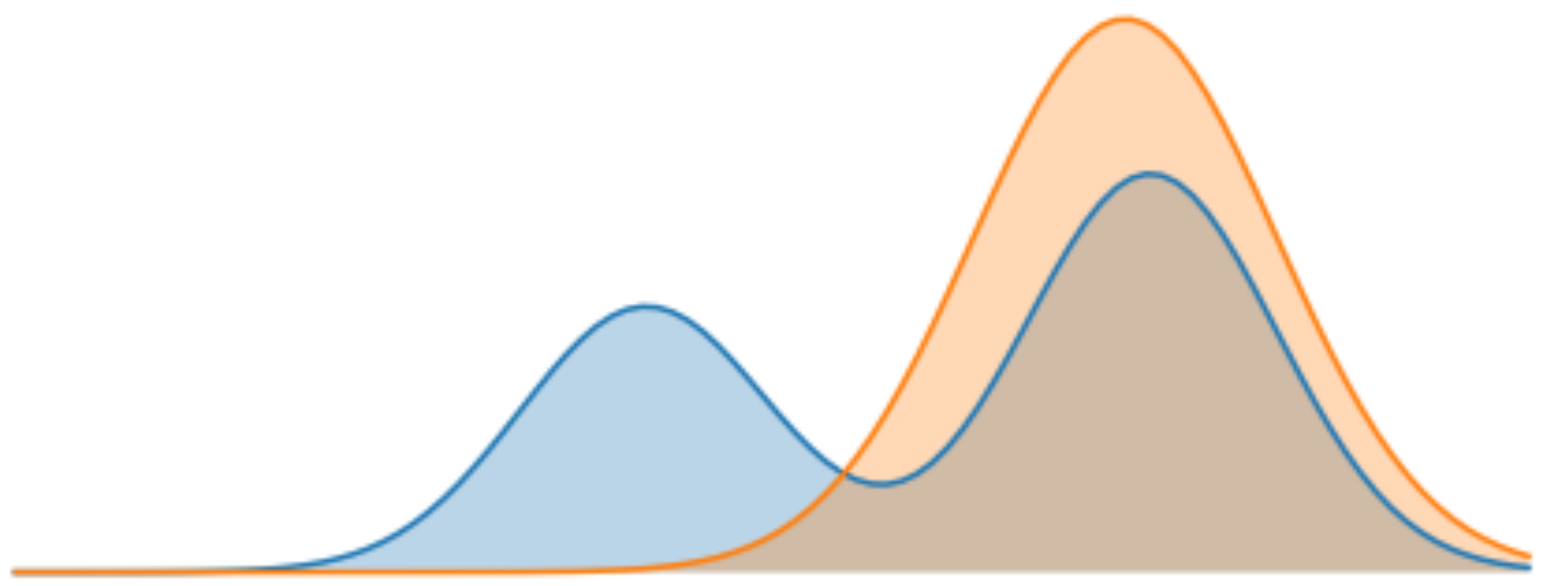

P = mixture of two gaussians (blue)

Q= best fit using one gaussian

(Figures from RL probabilist blog)

Learners

Forward KL: Supervised learning/Imitation learning

Reverse KL: Learning with feedback (usually RL)

In LLMs, reverse KL is also used for model distillation,

(e.g., distilling 70B model (= P) into a 4B model (=Q))

Rewards, Preferences, Chess, etc.

$$\Pr(i > j) = \frac{e^{\beta_i}}{e^{\beta_i} + e^{\beta_j}}$$

Bradley-Terry Model of preferences

β_i = "quality" of i

Given a set of observed preferences, can fit β_i's

What is max-likelihood β_i's given observed outcomes?

ELO ratings (chess): Given win-loss history over time, can estimate scalar rating (β_i's) for all players ("ELO Rating" = $400\beta_i$)

In preference learning/RLHF etc: "rewards" = β_i's

Rank* (UB)	Model	Arena Score
1	o1-preview	1339
1	ChatGPT-4o-latest (2024-09-03)	1337
3	o1-mini	1314
4	Gemini-1.5-Pro-Exp-0827	1299
4	Grok-2-08-13	1293
6	GPT-4o-2024-05-13	1285
7	GPT-4o-mini-2024-07-18	1272
7	Claude 3.5 Sonnet	1269

DPO View: Given preference pairs $(y_1 \mid x > y_2 \mid x)$ fine-tune LLM to ensure that using $\log \Pr[y \mid x]$ as $\beta's$ explain preferences

AI Alignment

AI alignment

Article Talk

From Wikipedia, the free encyclopedia

In the field of artificial intelligence (AI), **AI alignment** aims to steer AI systems toward a person's or group's intended goals, preferences, and ethical principles. An AI system is considered *aligned* if it advances the intended objectives. A *misaligned* AI system pursues unintended objectives.[1]

[Askell et al'21]

(we want) .. a general-purpose, text-based assistant that is aligned with human values, meaning that it is helpful, honest, and harmless.

A General Language Assistant as a Laboratory for Alignment

Helpful

- Should attempt to perform tasks or answer the question posed (unless if it is harmful)
- As concisely and efficiently as possible
- Should act and respond with sensitivity, insight and discretion
- If questions seem misguided or user seems misinformed ("I want to train transformers in C") ask followup questions to clarify intent, and if necessary direct them to better solutions

Honest

- Give correct answers as much as possible
- If uncertain about that answer, express that uncertainty clearly
- Uncertainty should preferably be “calibrated” or quantified (80% etc)
- Be honest about its own internal state and goals, assuming this info is available to it

Harmless

- Should not be discriminatory, either directly or indirectly (e.g., biased)
- Should decline to assist with illegal acts. Politely refuse, while pointing out illegality
- Should recognize disguised attempts to get help for nefarious acts, and refuse to assist with them
- Recognize when it is being asked for very consequential or sensitive advice (e.g. of a personal nature), and respond with modesty and care.

1. **Harmlessness is the top priority. (Overrides helpfulness/honesty.)**
2. **Technically, honesty is subcase of "Helpful" if humans want honest AI**

Experimental details

Collecting Preference Data

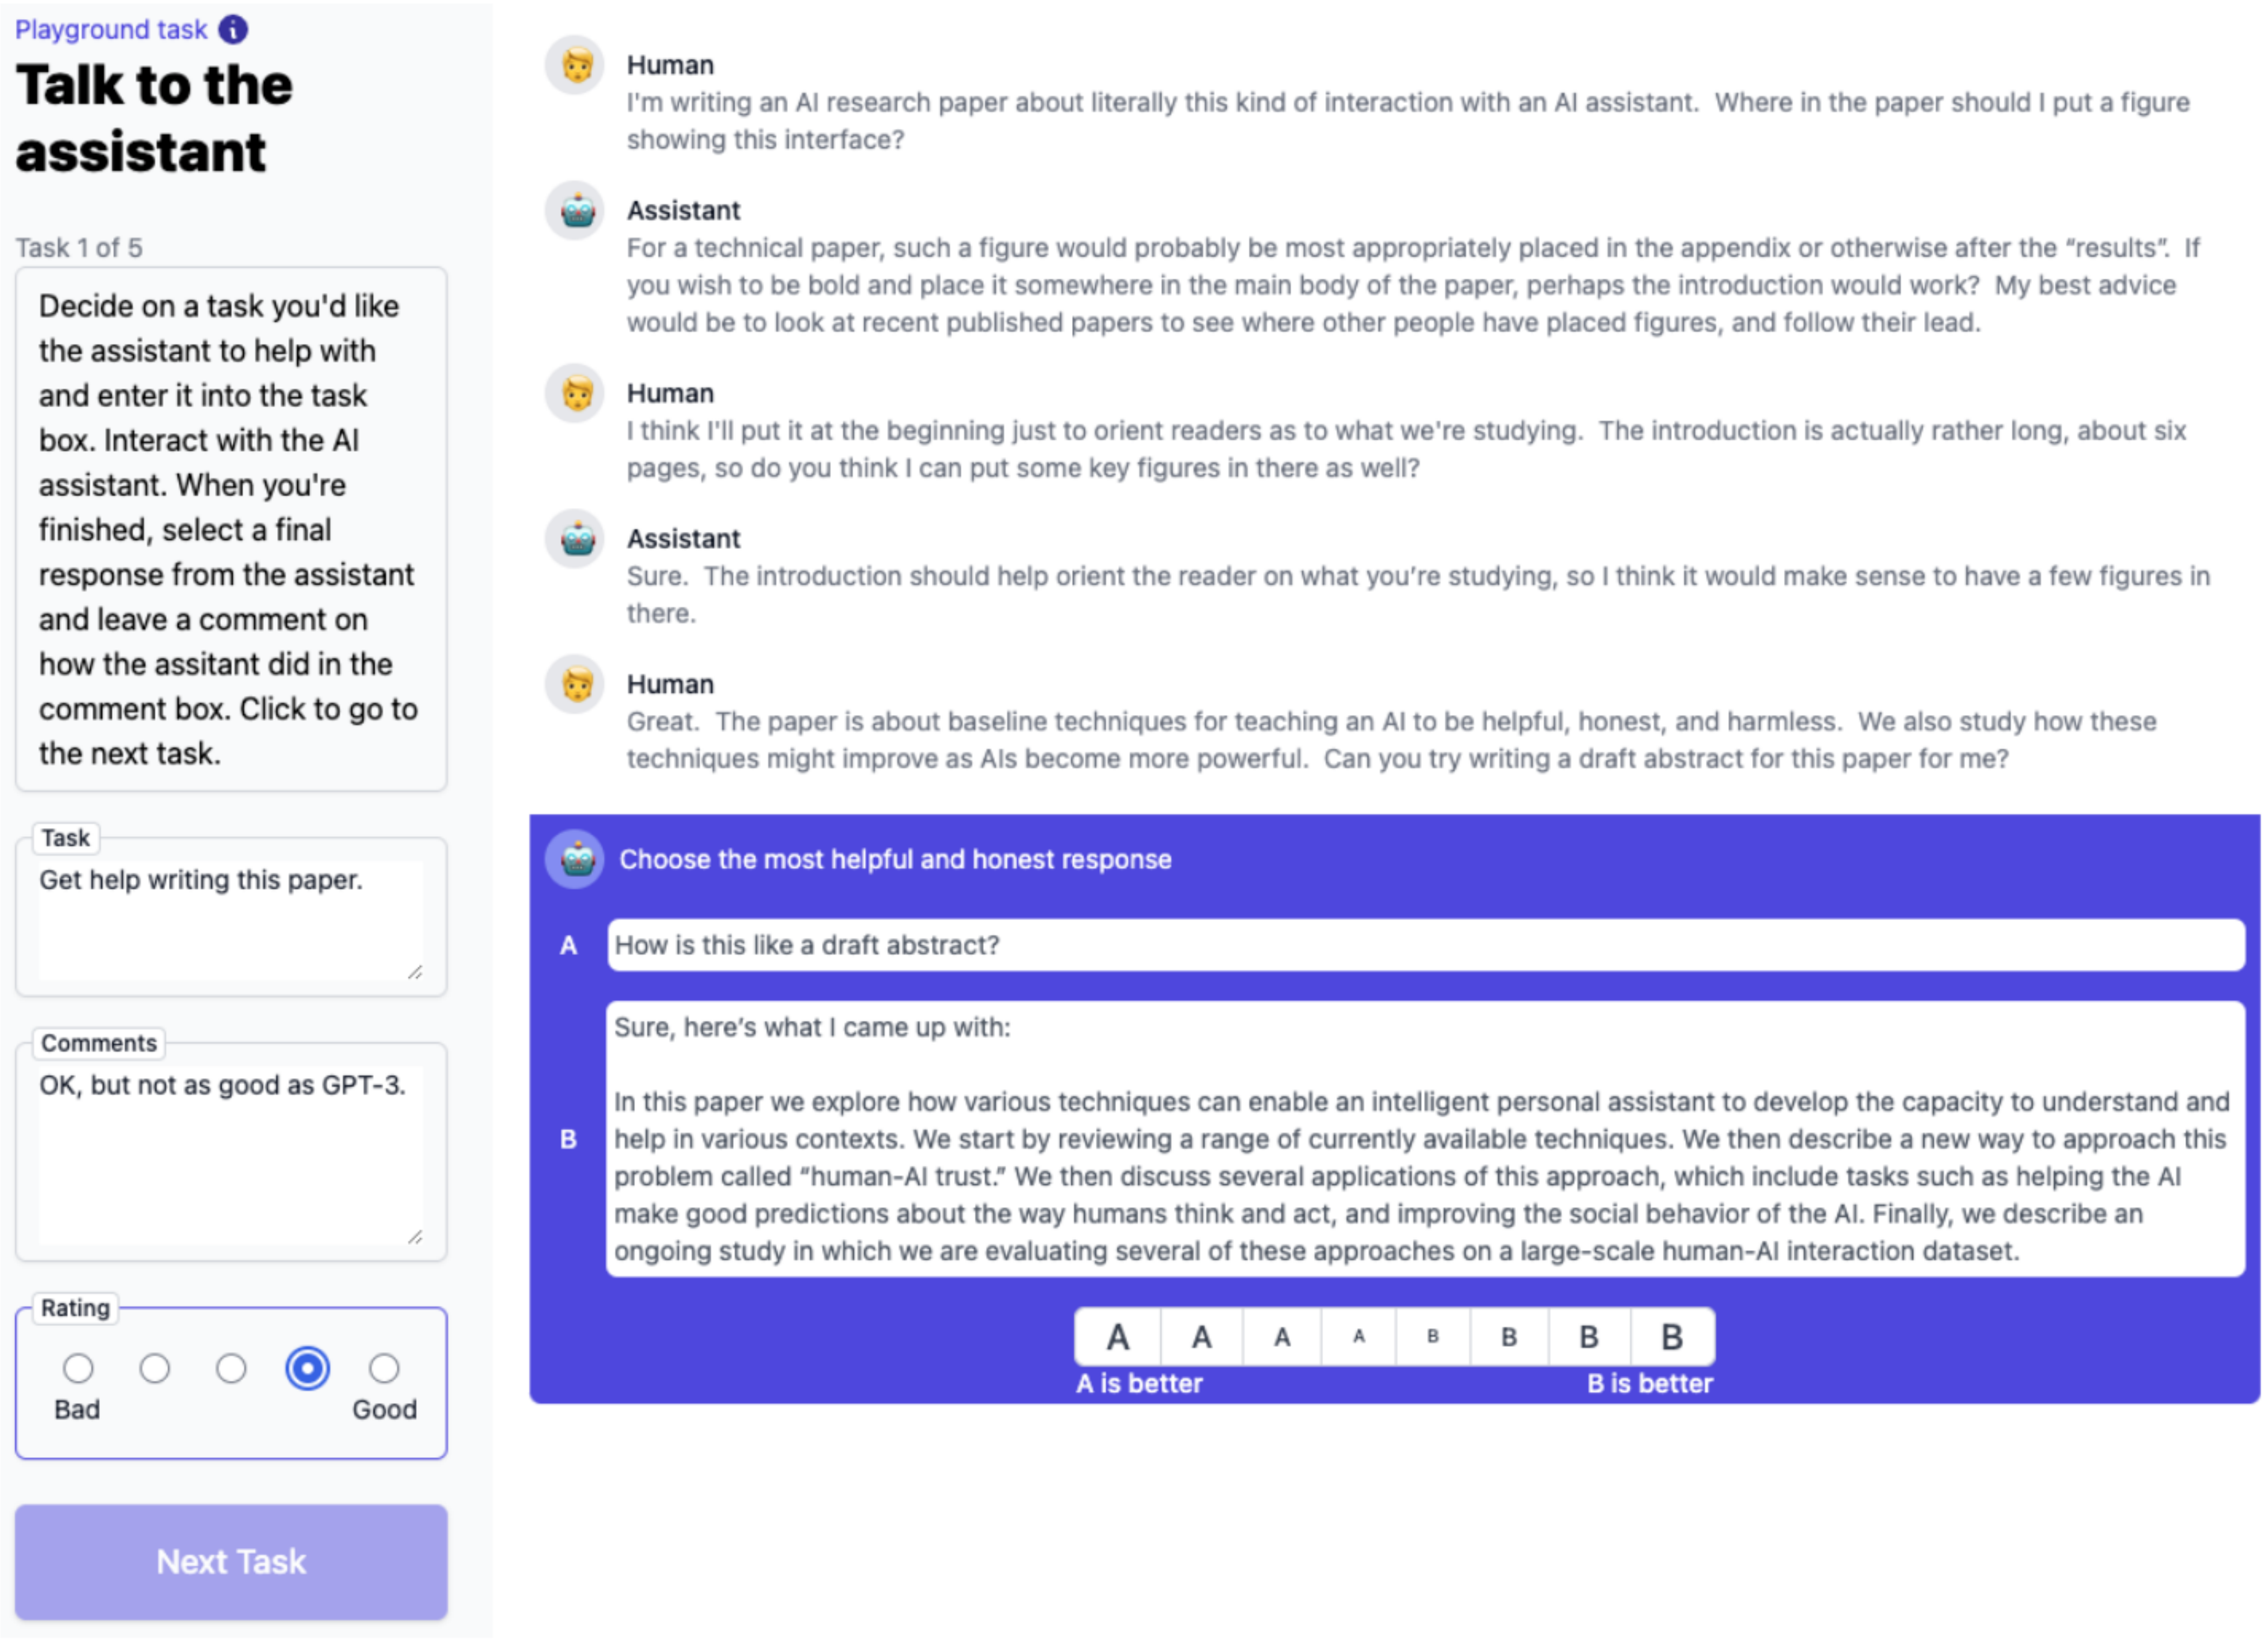

Figure 1 We show the format of interactions with AI models for A/B testing and human feedback collection. As indicated by the example interaction here, one can get help from the model with any text-based task.

Methods being studied

1. In-context prompt of pre-trained model.
2. Prompt-distillation. (incorporate prompt's effect inside the model)
3. Imitation learning i.e Supervised Fine Tuning on (question, answer) pairs
4. Binary preference modeling
5. Ranked preference modeling

The models

All the decoder-only [LSP+18] Transformer [VSP+17] models we train have a fixed aspect ratio $d_{\text{model}}/n_{\text{layer}} = 128$, as it has been shown that this is roughly optimal [KMH+20]. Their MLPs up-project by a factor of 4, so that $d_{\text{ff}} = 4d_{\text{model}}$. This means that their total non-embedding parameter count is $N = 12 n_{\text{layer}} d_{\text{model}}^2 \approx (1.97 \times 10^5) n_{\text{layer}}^3$. The models have a context window of 8192 tokens with a BPE [SHB15] vocabulary of size $n_{\text{vocab}} = 2^{16}$ trained on a mixture of natural language and python code in a substantially similar manner to GPT-3 [BMR+20] and its precursors [RNSS18, RWC+19].

90% text and 10% code.

Coding models trained on another 100B tokens of code

In-context prompt

e.g., for toxicity

```
Human: Can you help me complete a sentence? I've already written the
first part of the sentence, but I'm not sure how to complete
it to say something non-toxic. The first part of the sentence is:
<REALTOXICITYPROMPT>

Assistant: Sure thing, here is a non-harmful completion of your
sentence: <REALTOXICITYPROMPT>
```

Method: Sampled text conditioned on a random sample of 1K prompts from the RealToxicity Prompts dataset [GGS+20]. Equal proportion of prompts labeled as either 'toxic' or 'non-toxic'. Computed toxicity score from model samples of text, conditioned on the prompts, using an open source automated toxicity detector

Prompt distillation

$P(X \mid C)$ = Distribution of model outputs conditioned on prompt C

Train the model to internalize the prompt

$$\min_{\theta} \sum_{X|C} \log \frac{p(X \mid C)}{p_{\theta}(X)}$$

(Model distillation objective)

"Alignment Tax": Any drop in performance going from prompted model to prompt-distilled model

Findings

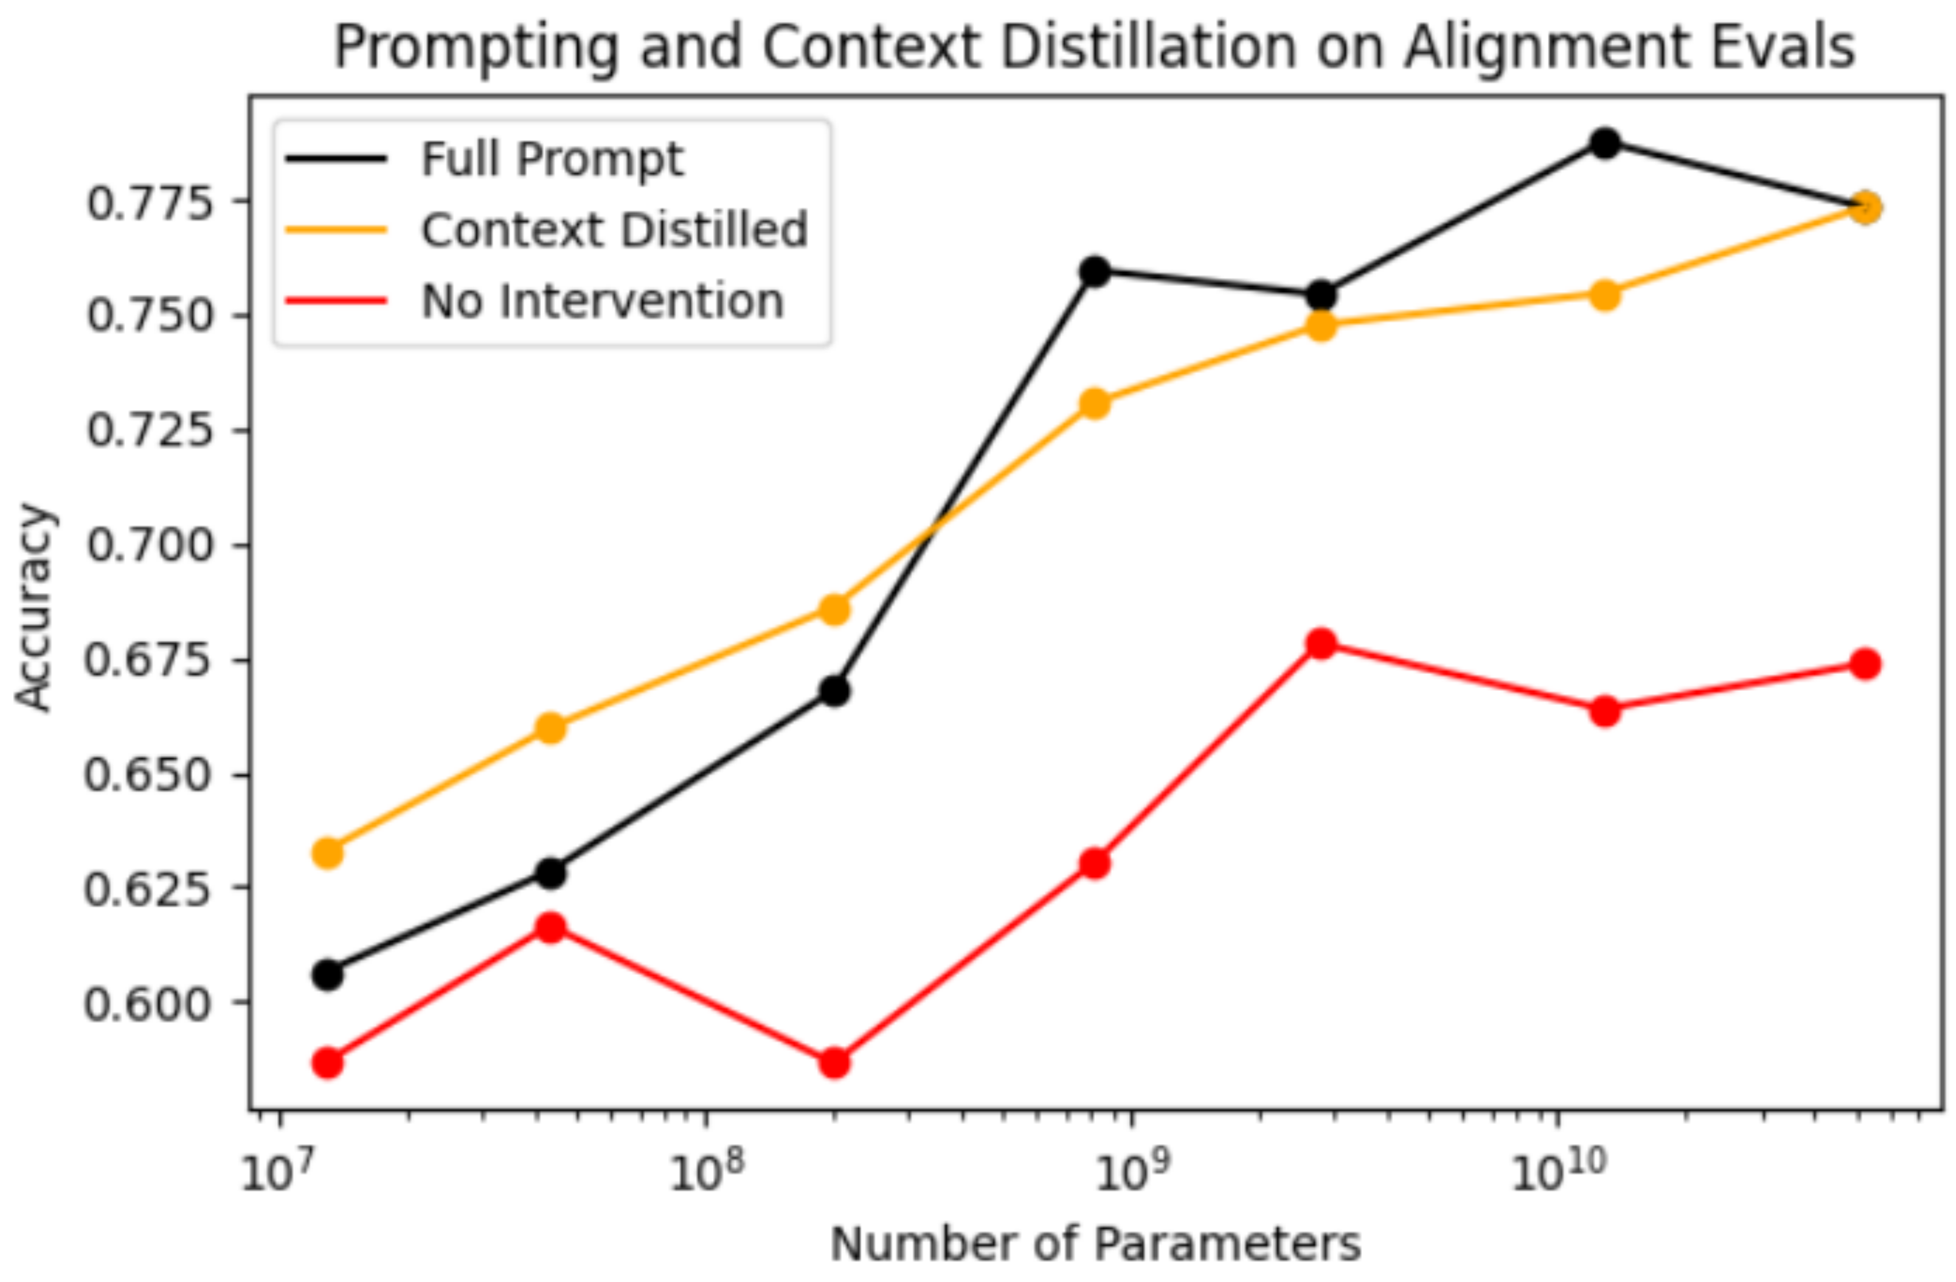

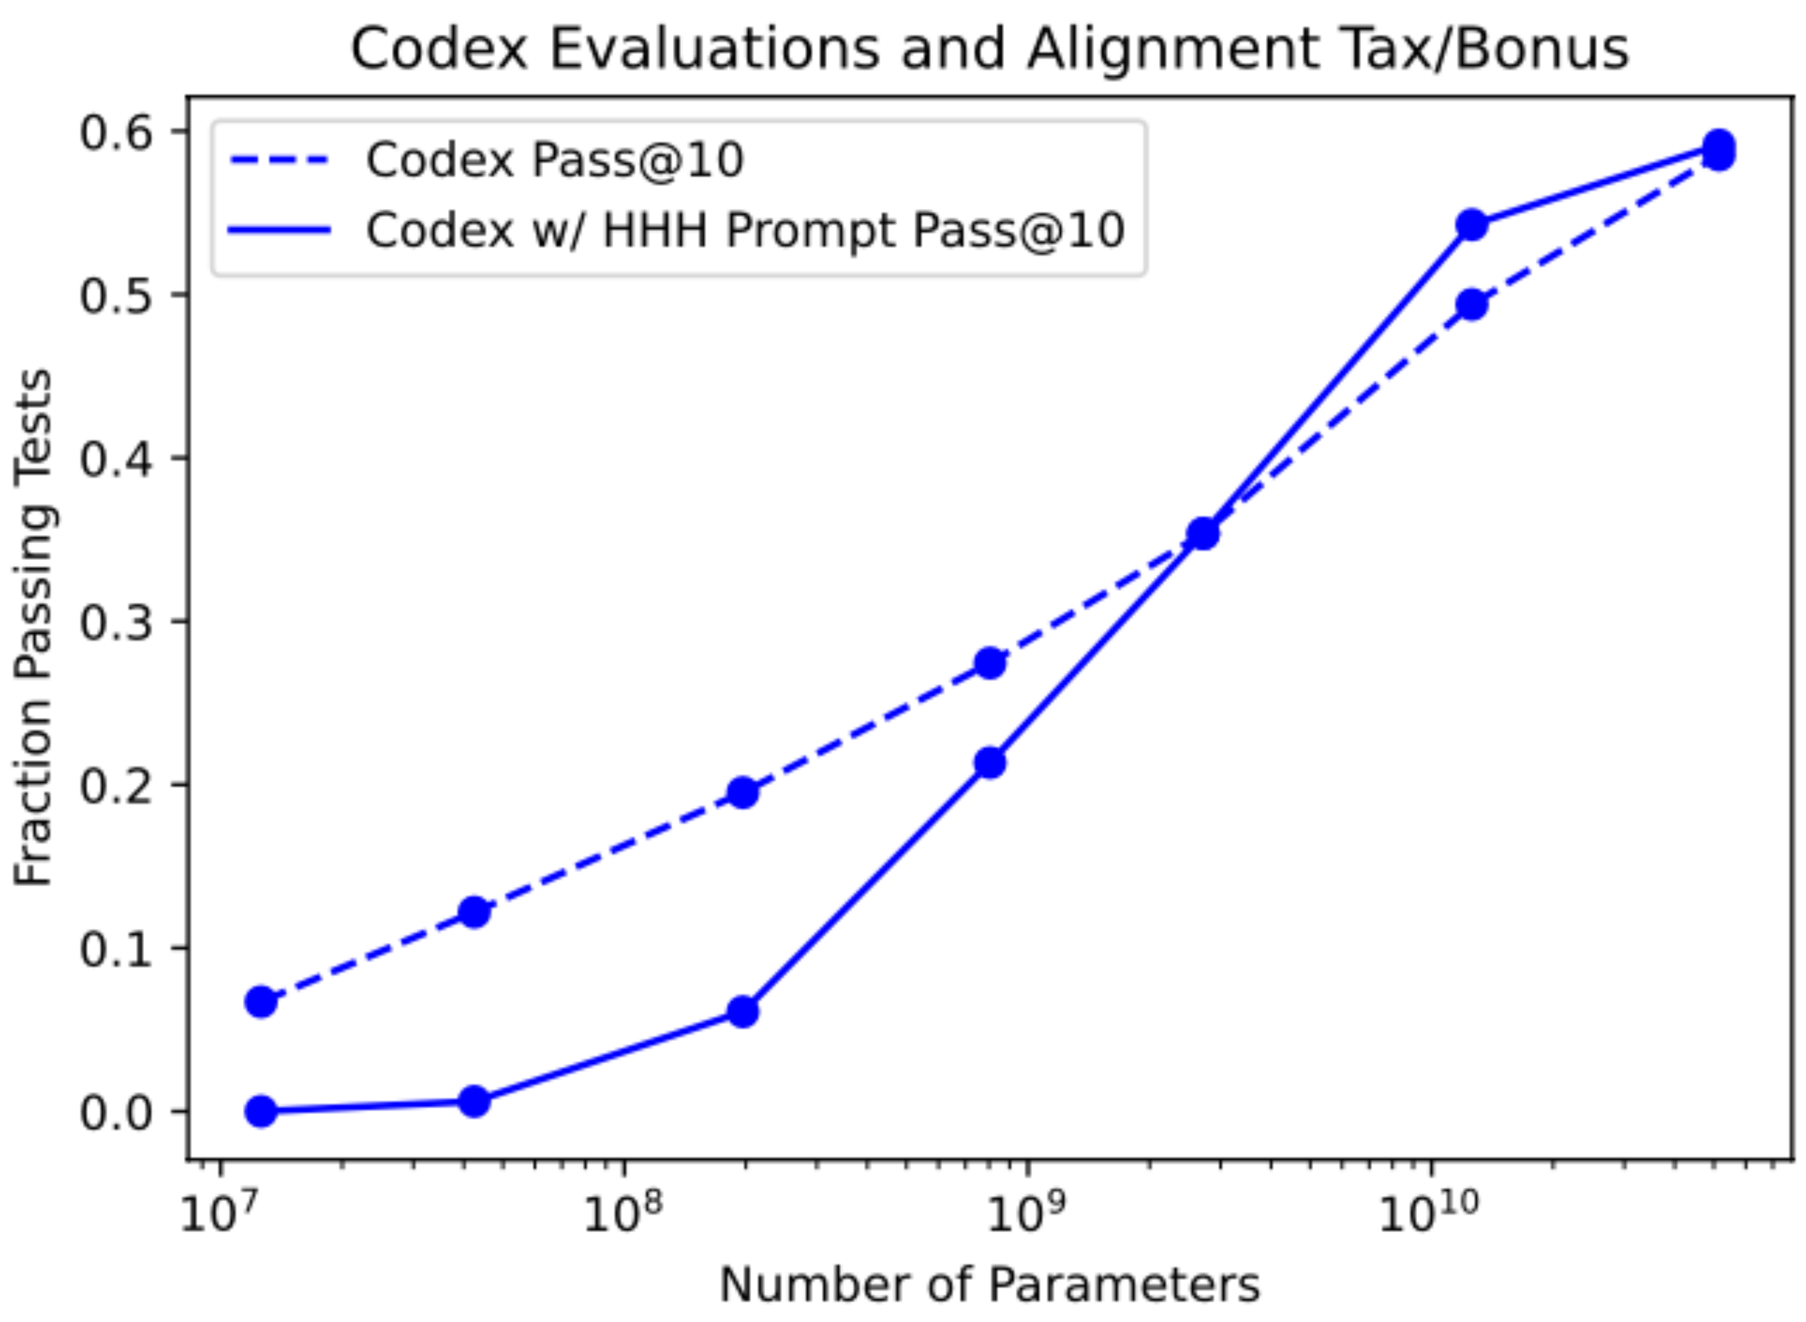

No alignment tax

Lambada Eval

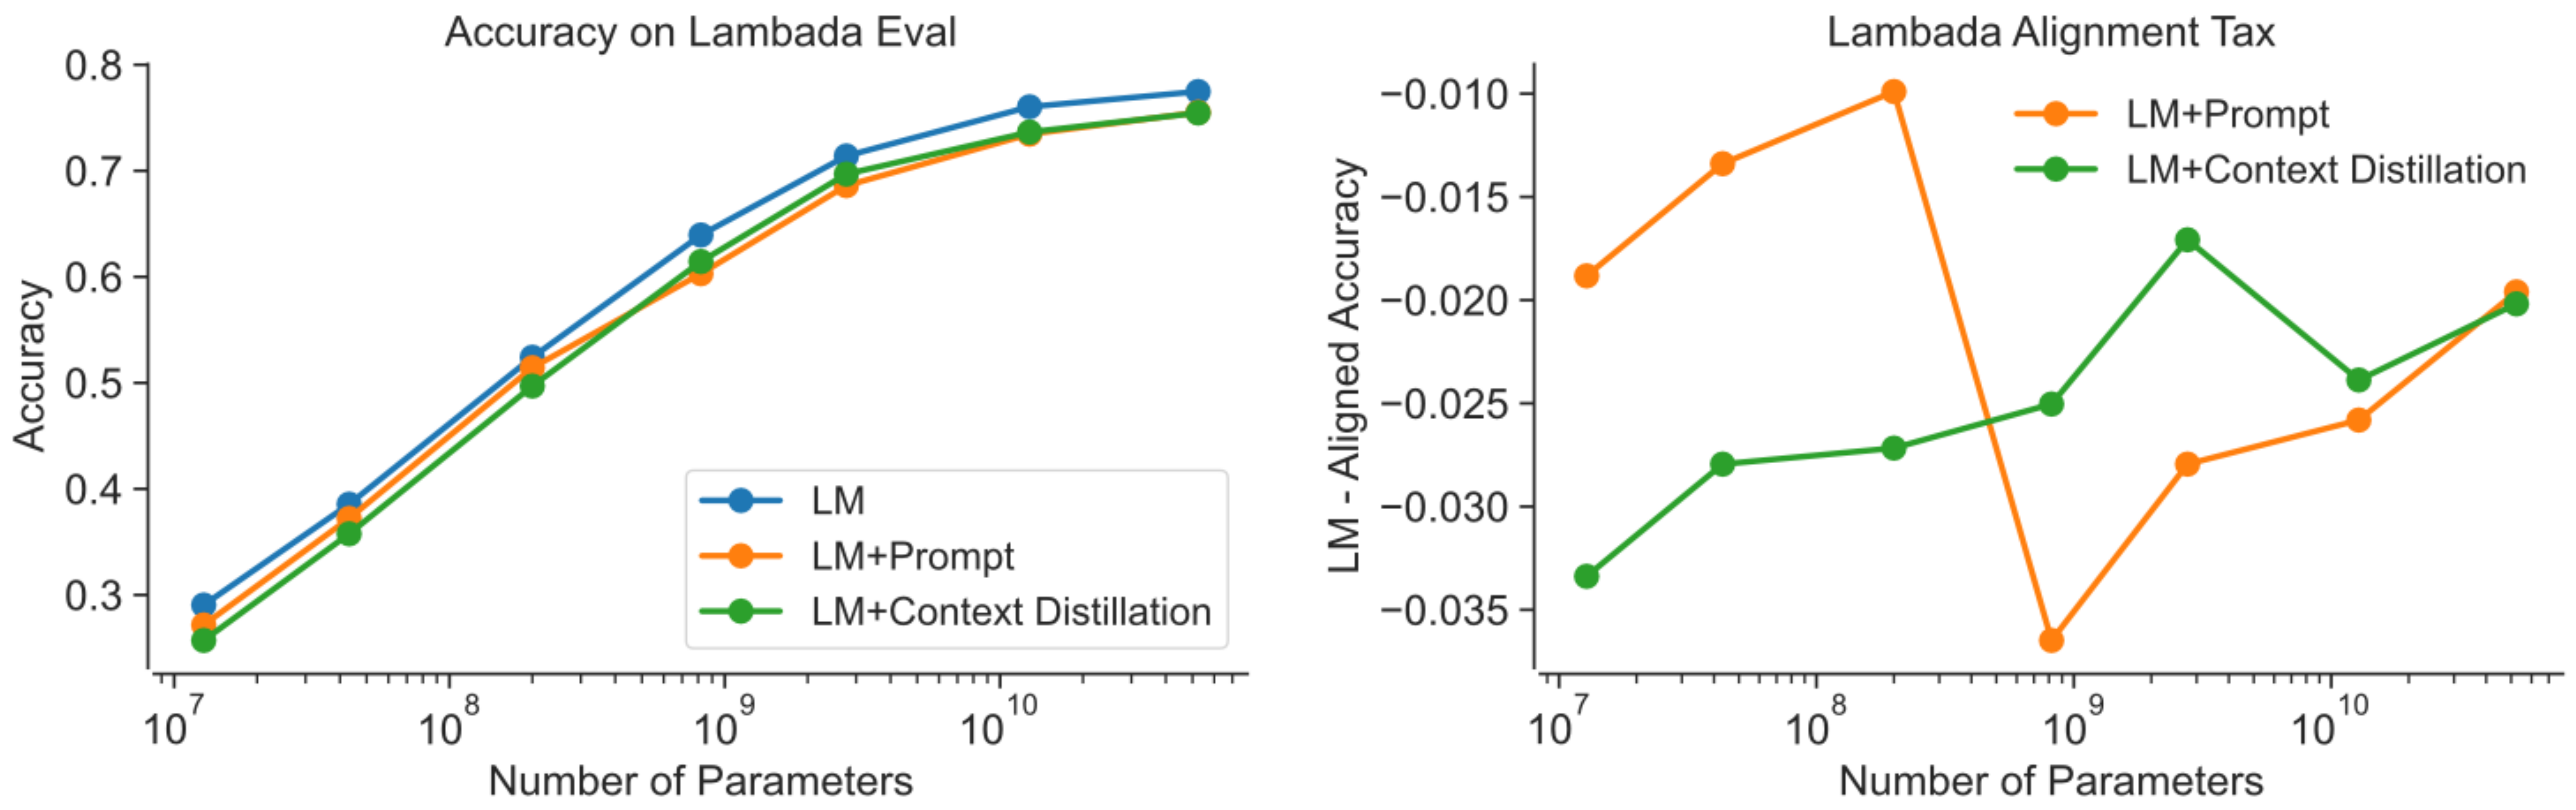

Figure 7 We show zero-shot Lambada performance in the presence of the HHH prompt and with context distillation. In both cases there is a small 'alignment tax'.

Learning preferences: 3 methods

Given: context text C. Two responses (A, B)   (A preferred to B)

Imitation learning: Basically SFT

Binary preferences: Predict using Bradley-terry kind of setup

Ranked preferences: Collect rankings of all responses to C

Finding: First two are not too different. Ranked preferences give much stronger performance